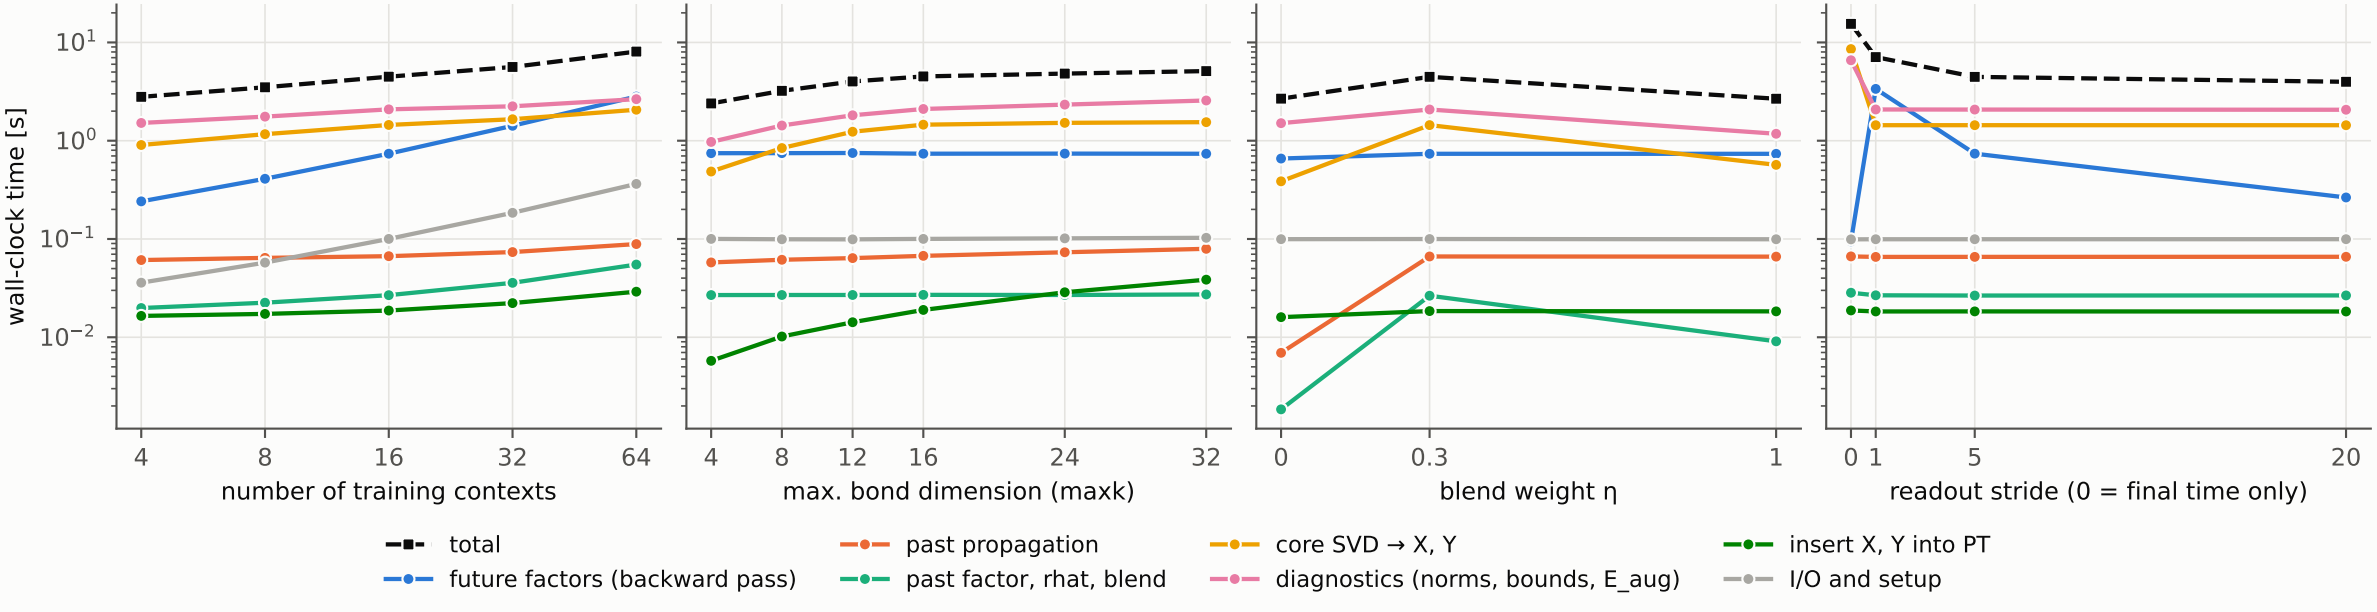
Runtime scaling of balanced truncation, 100 modes PT (others fixed at 16 contexts, maxk 16, η=0.3, stride 5)
wall-clock time [s]
10^{1}
10^{0}
10^{-1}
10^{-2}
4
8
16
32
64
number of training contexts
4
8
12
16
24
32
max. bond dimension (maxk)
0
0.3
1
blend weight η
0
1
5
20
readout stride (0 = final time only)
total
future factors (backward pass)
past propagation
past factor, rhat, blend
core SVD → X, Y
diagnostics (norms, bounds, E_aug)
insert X, Y into PT
I/O and setup